



PLAYERUNKNOWN'S  
**BATTLEGROUNDS**

Rachana Patil

# AGENDA

 Overview of the dataset

 Problem Statement

 Literature Review

 Data cleaning and preprocessing

 Method implementations – MLR, CART, kNN

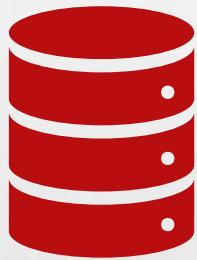
 Analysis

 Conclusion

# INTRODUCTION

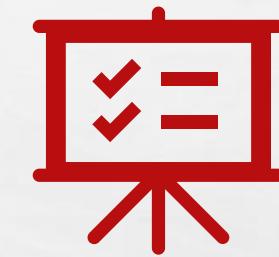
- GAME BACKGROUND
- SOURCE: [HTTPS://WWW.KAGGLE.COM/C/PUBG-FINISH-PLACEMENT-PREDICTION](https://www.kaggle.com/c/pubg-finish-placement-prediction)

# OVERVIEW OF THE DATASET



**variables and records in the data**

29 Variables : 4.4 million records



**Types of variables**

28 continuous, 1 categorical

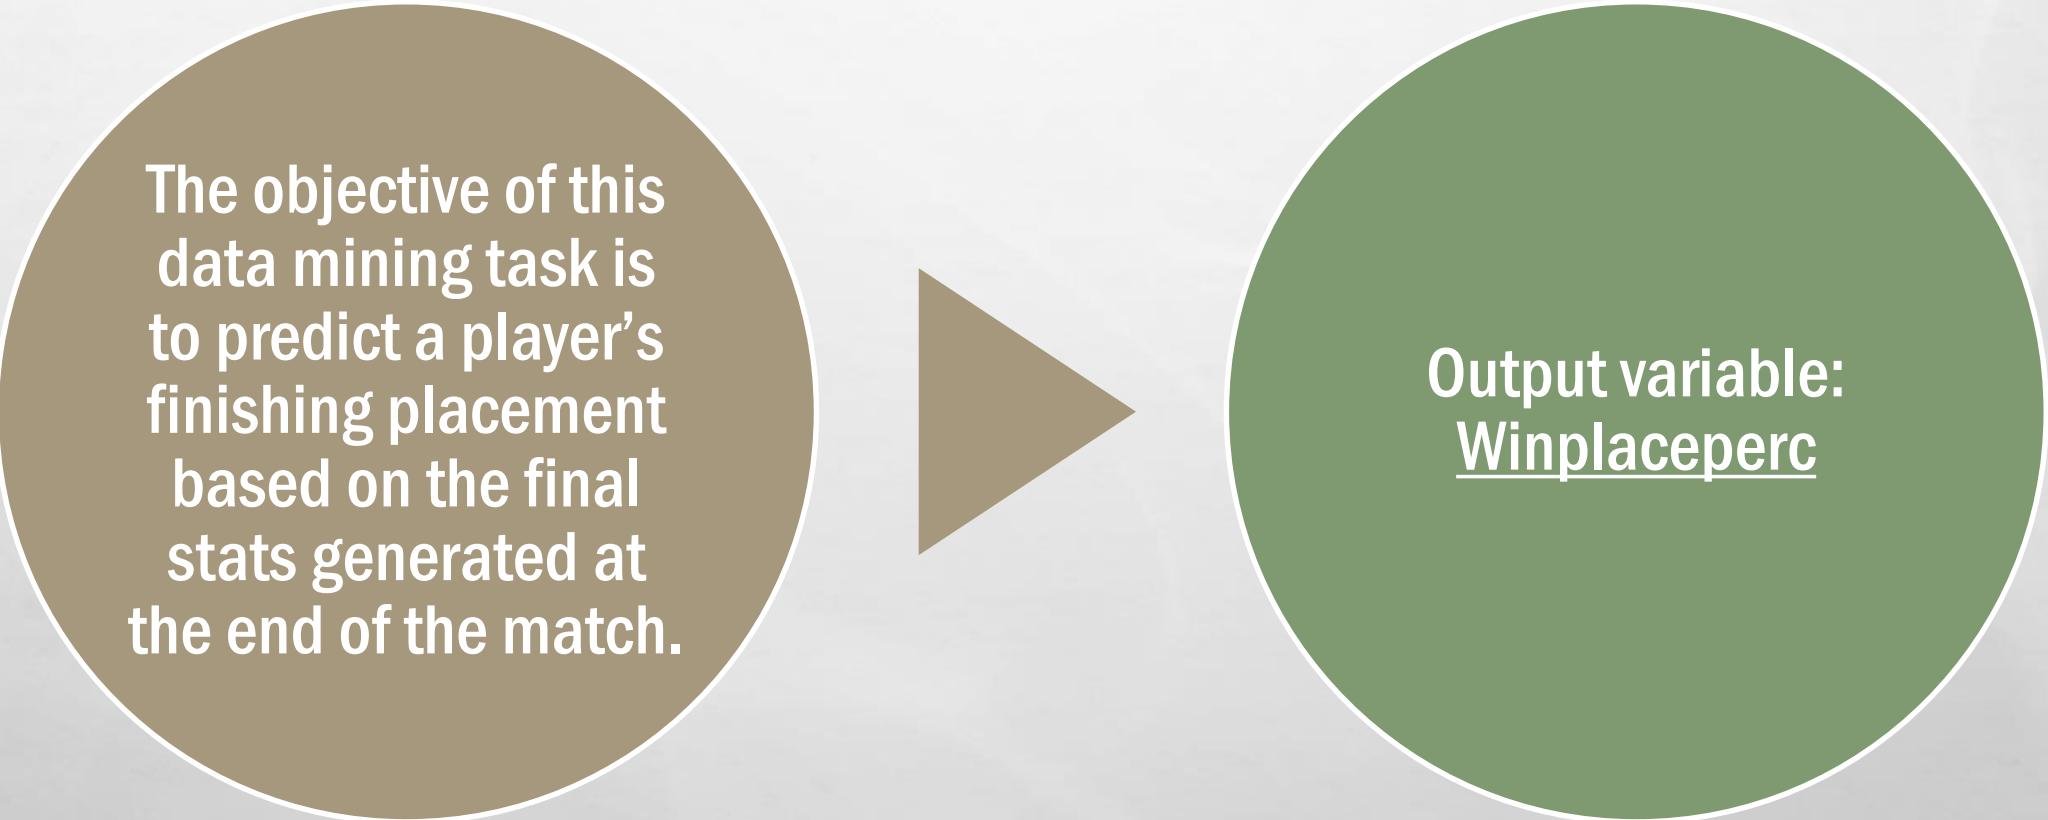
# OVERVIEW OF THE DATASET (VARIABLES)

DBNOs	killPoints
assists	killStreaks
boosts	kills
damageDealt	longestKill
headshotKills	matchDuration
heals	matchId
id	matchType
killPlace	rankPoints

# OVERVIEW OF THE DATASET (VARIABLES)

revives	winPoints
rideDistance	groupId
roadKills	numGroups
swimDistance	maxPlace
teamKills	winPlacPerc
vehicleDestroys	
walkDistance	
weaponsAcquired	

# PROBLEM STATEMENT



The objective of this data mining task is to predict a player's finishing placement based on the final stats generated at the end of the match.

Output variable:  
Winplaceperc

```
Chap6_MLR_complete.R scatterplots_4M.R Main_rcode.R scatterplots.R  
Source on Save Run Source  
sample  
In selection  
5 library(plyr)  
6 library(tidyr)  
7 library(ggplot2)  
8 library(gplots)  
9 library(rpart)  
10 library(rpart.plot)  
11 library(FNN)  
12 library(kknn)  
13 #library(vcd)  
14 setwd("F:/CSUF/Fall 2018/ISDS 574/Project/original dataset/")  
15 df <- read.csv("train_V2.csv", head = TRUE, stringsAsFactor = FALSE)  
16 set.seed(1)  
17 df <- df[sample(nrow(df), 100000), ]  
18 dim(df)  
19 colnames(df)  
20 <  
19:1 (Top Level)
```

```
Console Terminal  
~/ ~/  
type citation()  
  
Type 'demo()' for some demos, 'help()' for on-line help, or  
'help.start()' for an HTML browser interface to help.  
Type 'q()' to quit R.  
  
[workspace loaded from ~/.RData]  
  
> df <- read.csv("train_V2.csv", head = TRUE, stringsAsFactor = FALSE)  
> set.seed(1)  
> df <- df[sample(nrow(df), 100000), ]  
[1:100000, ]
```

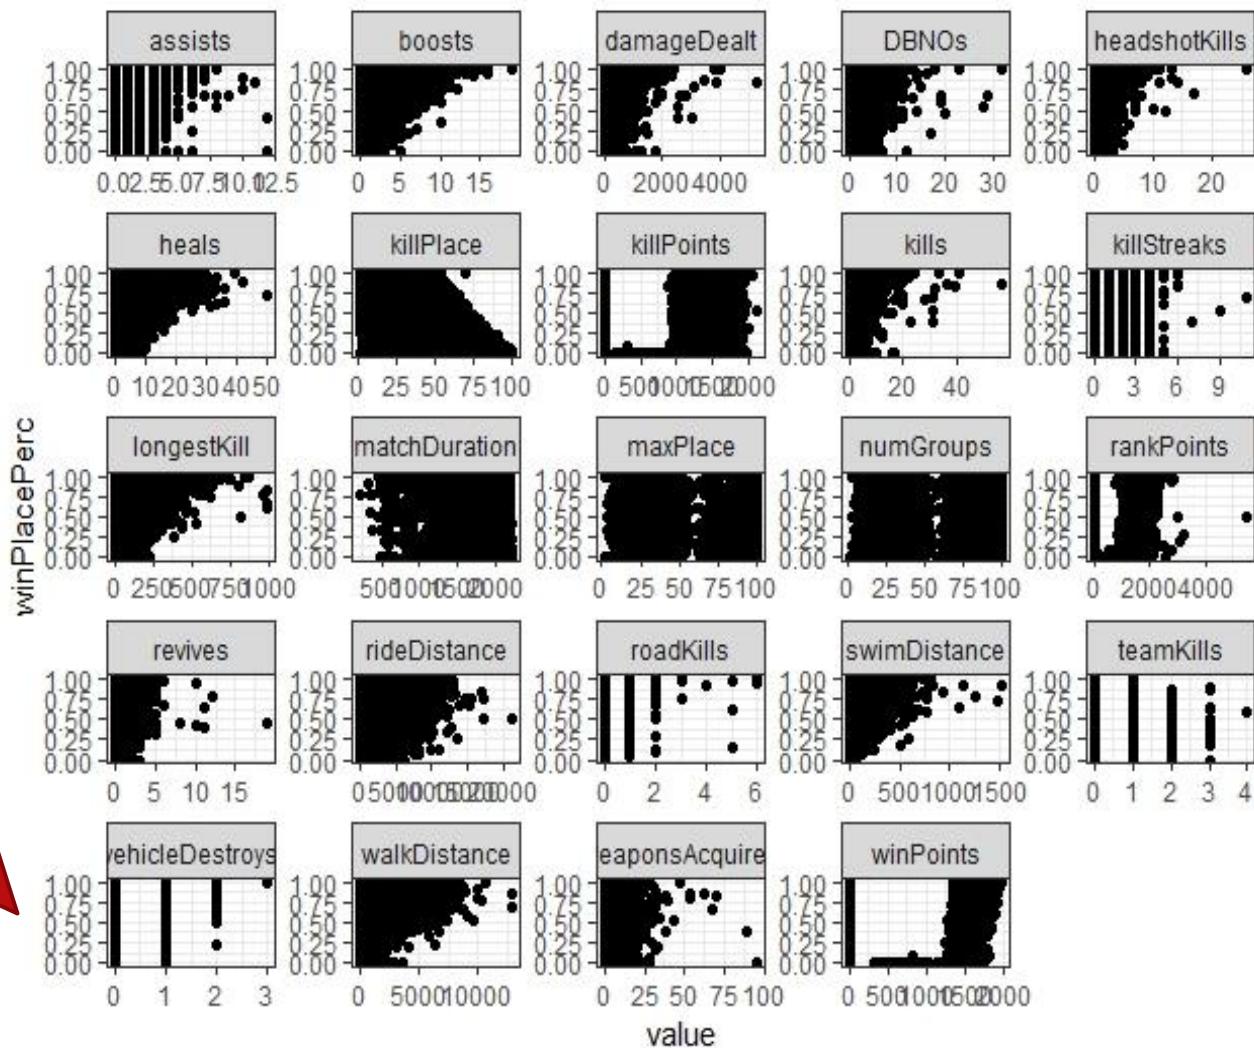


# DATA EXPLORATION

We had huge data set of 4.4 million rows of data which was difficult to handle. Thus we did the random sampling of 100,000 rows of data

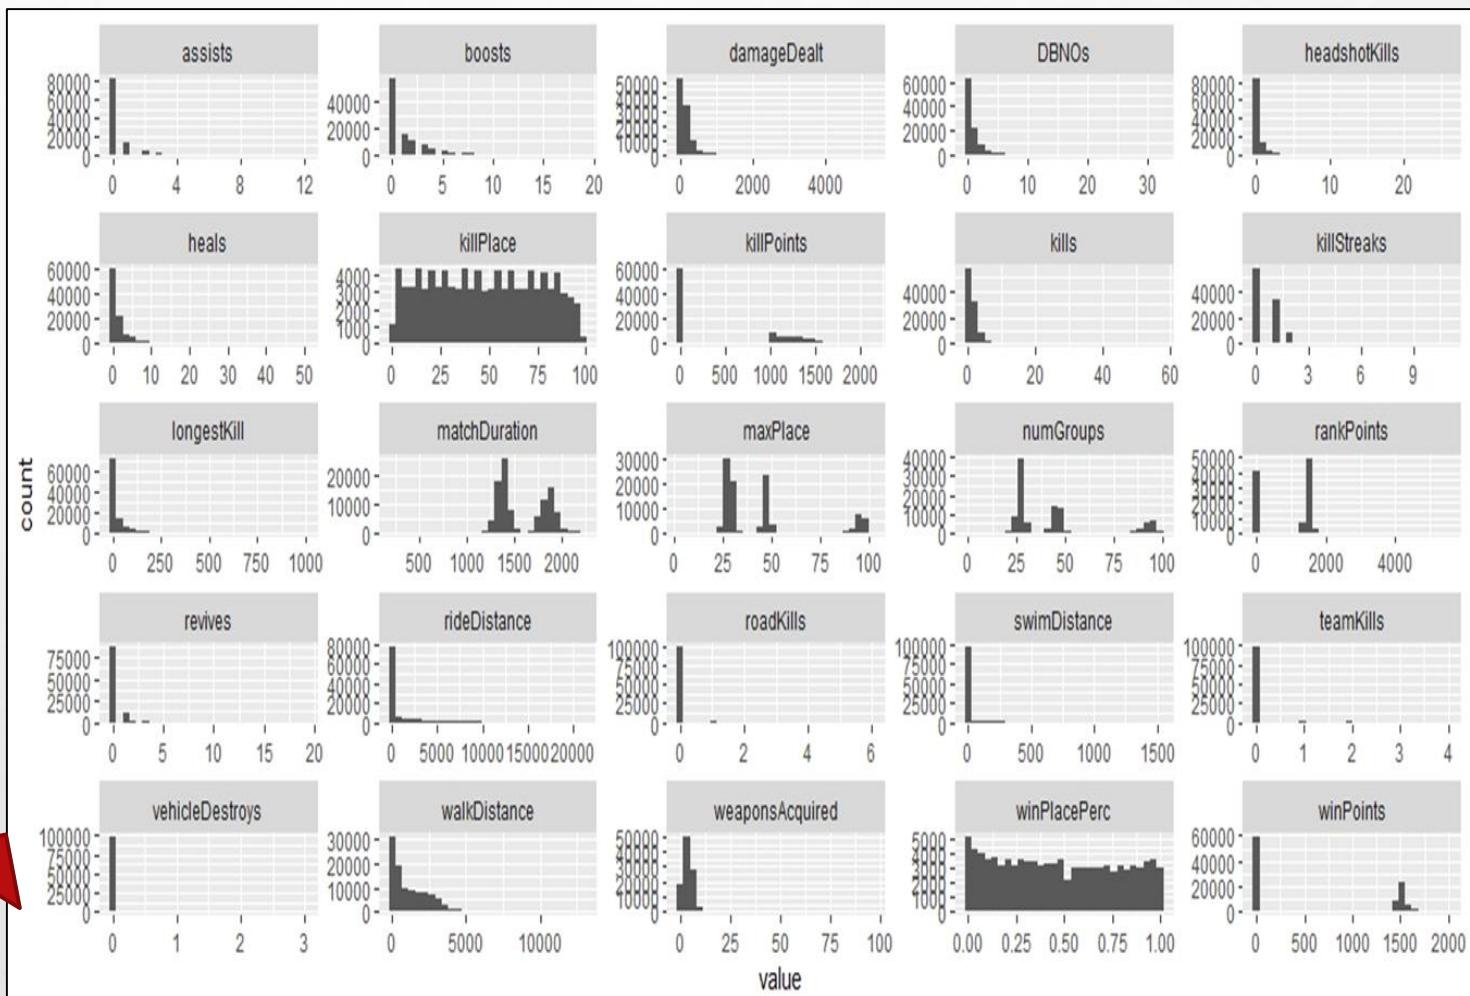
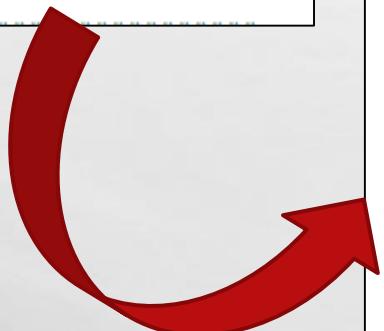
# SCATTERPLOTS

```
35  
36 ##### Consolidated Scatter Plot#####  
37 df %>%  
38 gather(-winPlacePerc, key = "var", value = "value") %>%  
39 ggplot(aes(x = value, y = winPlacePerc)) +  
40 geom_point() +  
41 facet_wrap(~ var, scales = "free") +  
42 theme_bw()  
43
```



# HISTOGRAMS

```
43  
44 ###### Consolidated Histogram #####  
45 df %>%  
46 keep(is.numeric) %>%  
47 gather() %>%  
48 ggplot(aes(value)) +  
49 facet_wrap(~ key, scales = "free") +  
50 geom_histogram()  
51
```



# DATA CLEANING AND PREPROCESSING

## Variable Reduction

Data Set	Rows, Columns
After Random Sampling	[100,000, 29]
Based On Domain Knowledge	[100,000, 24]

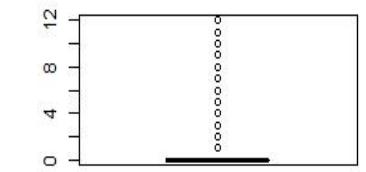
1



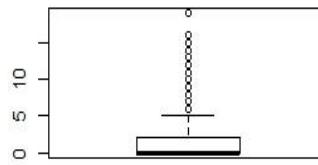
- 1) Group ID
- 2) Match ID
- 3) ID
- 4) Matchtype
- 5) RoadKills

# DATA CLEANING AND PREPROCESSING

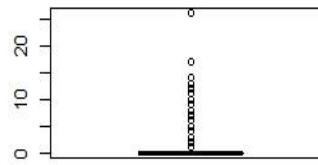
2



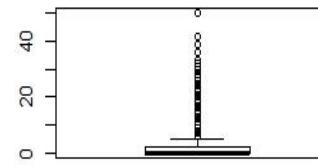
assists



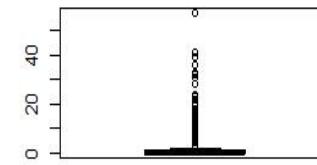
boosts



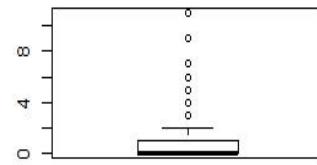
headshotKills



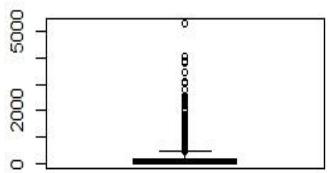
heals



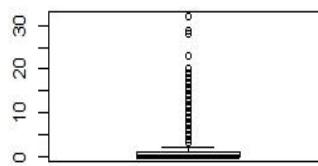
kills



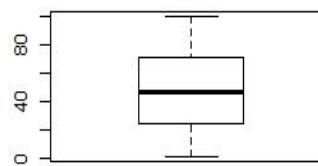
killStreaks



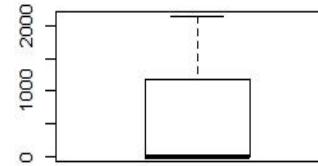
damageDealt



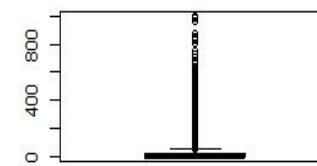
DBNOs



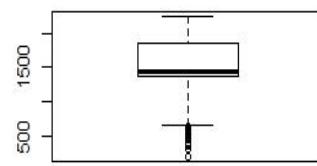
killPlace



killPoints



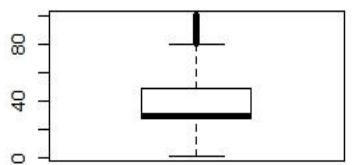
longestKill



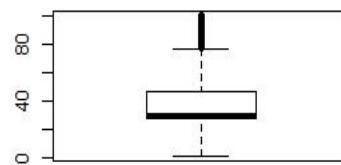
matchDuration

# DATA CLEANING AND PREPROCESSING

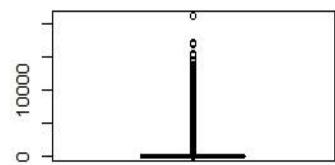
## Outlier Detection: Box Plot



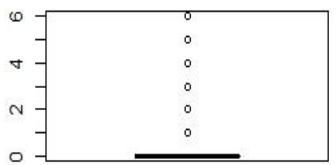
maxPlace



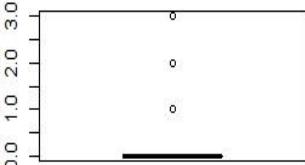
numGroups



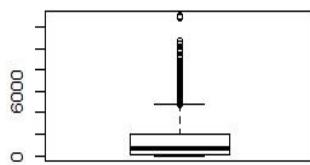
rideDistance



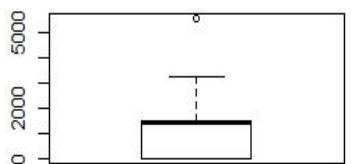
roadKills



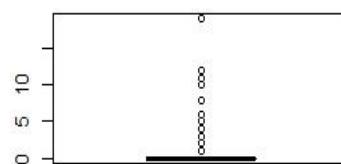
vehicleDestroys



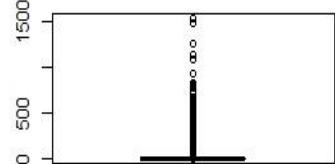
walkDistance



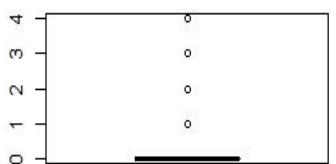
rankPoints



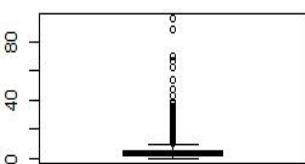
revives



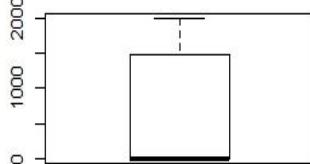
swimDistance



teamKills



weaponsAcquired



winPoints

# DATA CLEANING AND PREPROCESSING

3

# Outliers Handled

```
##### Data Cleaning #####
as <- which(df$assists > 1)
bo <- which(df$boosts >= 6)
Dma <- which(df$damageDealt >= 500)
DBno <- which(df$DBNOS >= 3)
hskill <- which(df$headshotkills >= 2)
```

[47044,24]

4

# Correlation Matrix

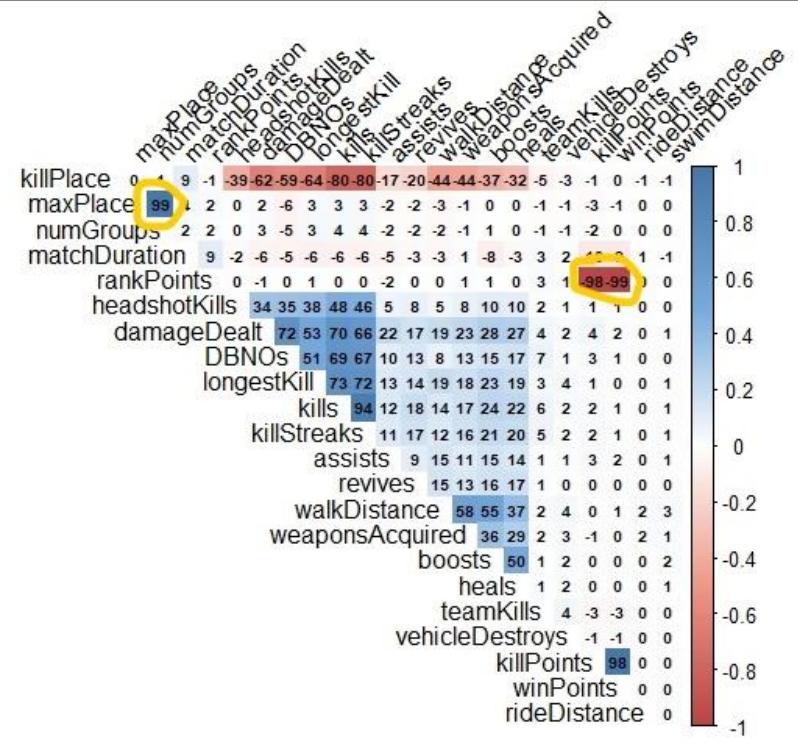
```
230 cormat<- cor(df3) #/// to create corelation matrix  
231  
232 corrplot(cormat,method = 'number', title = 'Corealation Matrix',  
233 addCoefasPercent = TRUE, number.cex = 0.65)  
234
```

## **Cutoff Value: 0.90**

MaxPlace, Numgroup	99
Killpoints, Rankpoints	-98
Winpoints, Rankpoints	-99

[47044,21]

3



# DATA CLEANING AND PREPROCESSING

## 6 Normalization

```
273  
274 df3 <- as.data.frame(scale(df3[,c(1:20)]))  
275  
276 class(df3)|  
277 summary(df3)
```



On all variables except Output variable

```
> summary(df3)  
  assists      boosts      damageDealt  
 Min. :-0.3387  Min. :-0.4252  Min. :-0.8609  
 1st Qu.:-0.3387 1st Qu.:-0.4252 1st Qu.:-0.8609  
 Median :-0.3387 Median :-0.4252 Median :-0.2882  
 Mean   : 0.0000  Mean   : 0.0000  Mean   : 0.0000  
 3rd Qu.:-0.3387 3rd Qu.:-0.4252 3rd Qu.: 0.3873  
 Max.   : 2.9522  Max.   : 5.9531  Max.   : 5.3053
```

## 7 Partitioning Data Set

70% Training Dataset

```
trainind <- sample(1:nrow(df3), nrow(df3)*.7)  
train_df <- df3[trainind, ]  
length(trainind)
```

30% Test Dataset

```
testind <- setdiff(1:nrow(df3), trainind)  
test_df <- df3[testind, ]  
length(testind)
```

```
> dim(train_df)  
[1] 32930 21
```

```
> dim(test_df)  
[1] 14114 21
```

# METHOD IMPLEMENTATIONS

MLR – Multiple  
Linear Regression

CART –  
Classification and  
Regression Tree

kNN – k Nearest  
Neighbours

# MULTIPLE LINEAR REGRESSION

- Predict the value of the dependent variable `winPlacePerc` based on predictors `kills`, `DBNOs`, `assists`, `boosts` etc.
- Regression coefficients  $B_1, B_2, \dots, B_p$  in the equation and they are calculated using ordinary least squares (OLS).

$$Y = B_0 + B_1 X_1 + B_2 X_2 + \dots + B_p X_p + \text{ERROR}$$

# MULTIPLE LINEAR REGRESSION

## VARIABLE SELECTION METHODS

- FORWARD SELECTION

```
obj.null <- lm(winPlacePerc ~ 1, dat = train_df)
obj.full <- lm(winPlacePerc ~ ., dat = train_df)
obj_fwd = step(obj.null, scope=list(lower=obj.null, upper=obj.full), direction='forward')
summary(obj_fwd)
```

LM(FORMULA = WINPLACEPERC ~ WALKDISTANCE + KILLPLACE + KILLSTREAKS + NUMGROUPS + MATCHDURATION + KILLS  
+ BOOSTS + WEAPONSACQUIRED + ASSISTS + DBNOS + HEALS + REVIVES + KILLPOINTS + LONGESTKILL + HEADSHOTKILLS  
+ DAMAGEDEALT + SWIMDISTANCE, DATA = TRAIN\_DF)

```
Residual standard error: 0.1165 on 32912 degrees of freedom
Multiple R-squared:  0.7891,    Adjusted R-squared:  0.7889
F-statistic:  7242 on 17 and 32912 DF,  p-value: < 2.2e-16
```

# MULTIPLE LINEAR REGRESSION

## VARIABLE SELECTION METHODS

- BACKWARD SELECTION

```
obj.null <- lm(winPlacePerc ~ 1, dat = train_df)
obj.full <- lm(winPlacePerc ~ ., dat = train_df)
obj_bkwd = step(obj.null, scope=list(lower=obj.null, upper=obj.full), direction='backward')
summary(obj_bkwd)
```

LM(FORMULA = WINPLACEPERC ~ ASSISTS + BOOSTS + DAMAGEDEALT + DBNOS + HEADSHOTKILLS + HEALS + KILLPLACE + KILLPOINTS + KILLS + KILLSTREAKS + LONGESTKILL + MATCHDURATION + NUMGROUPS + REVIVES + SWIMDISTANCE + WALKDISTANCE + WEAPONSACQUIRED, DATA = TRAIN\_DF)

```
Residual standard error: 0.1165 on 32912 degrees of freedom
Multiple R-squared:  0.7891,    Adjusted R-squared:  0.7889
F-statistic: 7242 on 17 and 32912 DF,  p-value: < 2.2e-16
```

# MULTIPLE LINEAR REGRESSION

## VARIABLE SELECTION METHODS

- STEPWISE SELECTION

```
obj.null <- lm(winPlacePerc ~ 1, dat = train_df)
obj.full <- lm(winPlacePerc ~ ., dat = train_df)
obj_step = step(obj.null, scope=list(lower=obj.null, upper=obj.full), direction='both')
summary(obj_step)
```

LM(FORMULA = WINPLACEPERC ~ WALKDISTANCE + KILLPLACE + KILLSTREAKS + NUMGROUPS + MATCHDURATION  
+ KILLS + BOOSTS + WEAPONSACQUIRED + ASSISTS + DBNOS + HEALS + REVIVES + KILLPOINTS + LONGESTKILL +  
HEADSHOTKILLS + DAMAGEDEALT + SWIMDISTANCE, DATA = TRAIN\_DF)

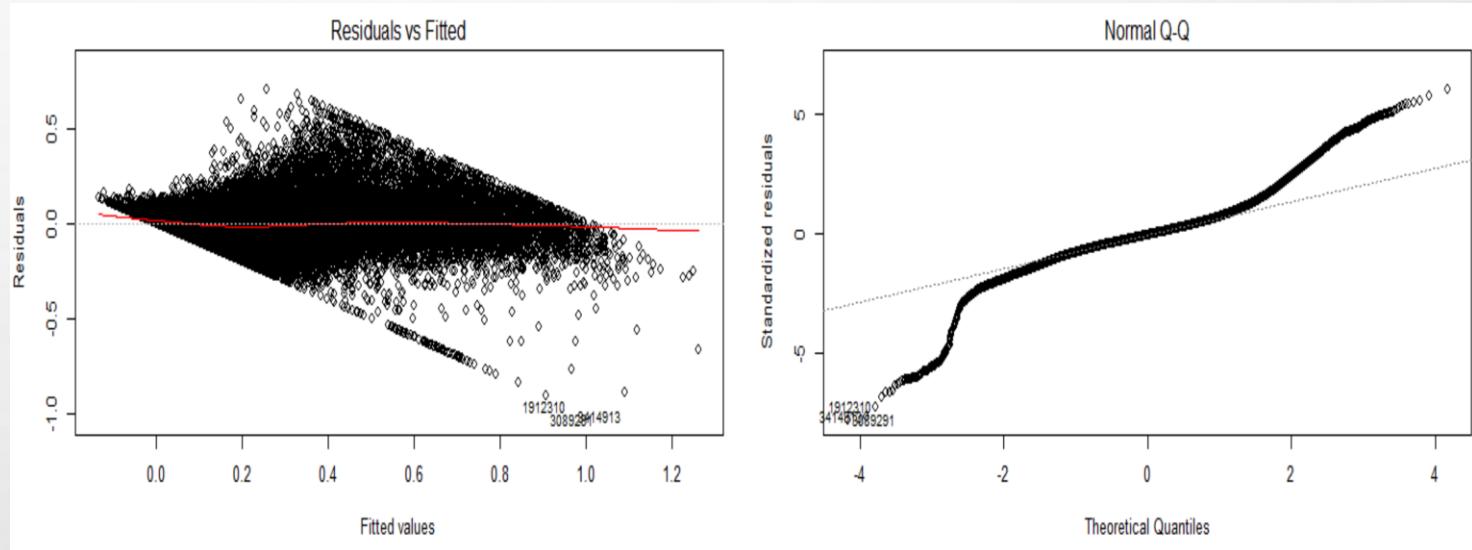
```
Residual standard error: 0.1165 on 32912 degrees of freedom
Multiple R-squared:  0.7891,    Adjusted R-squared:  0.7889
F-statistic:  7242 on 17 and 32912 DF,  p-value: < 2.2e-16
```

# MULTIPLE LINEAR REGRESSION

## ASSUMPTIONS

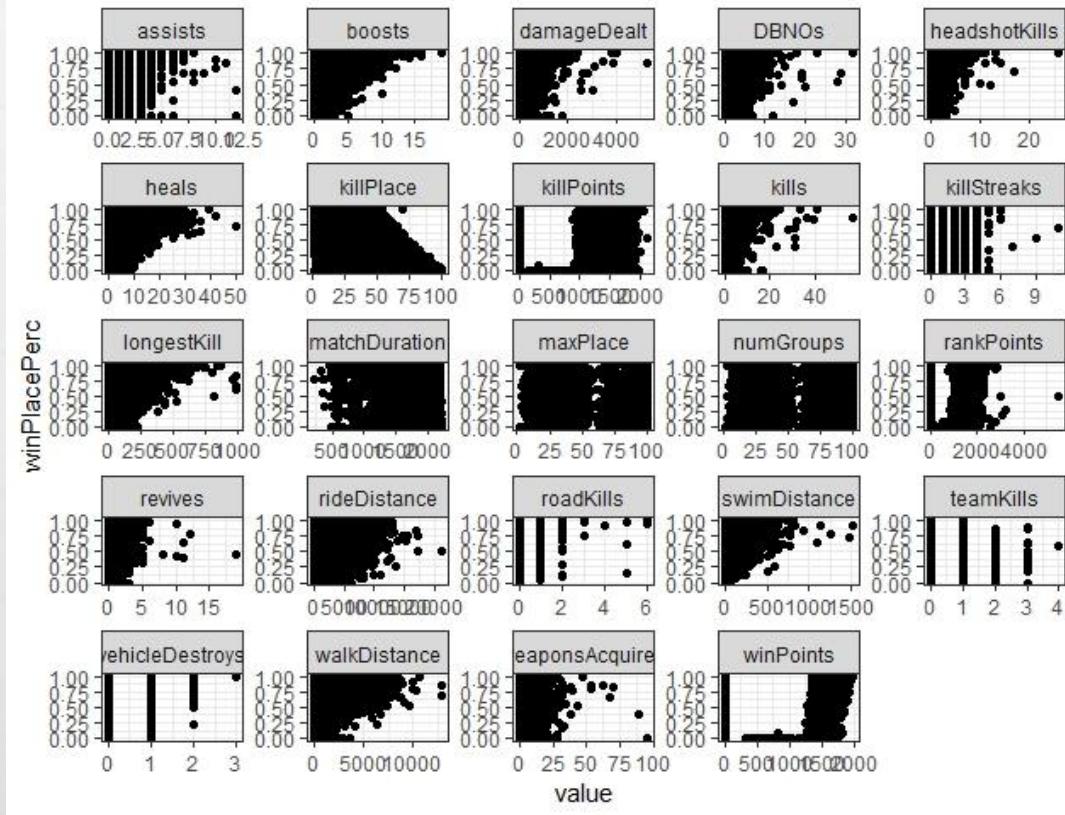
- **NORMALITY**

The residuals (errors) should be normally distributed about the predicted responses.



Two common methods to check this assumption include using: (a) a histogram (with a superimposed normal curve) and a Normal P-P Plot; or (b) a Normal Q-Q Plot of the studentized residuals.

# MULTIPLE LINEAR REGRESSION



## ASSUMPTIONS

- **LINEARITY**

Linear relationship between the outcome variable and the independent variables.

Scatterplots can show whether there is a linear or curvilinear relationship

# MULTIPLE LINEAR REGRESSION

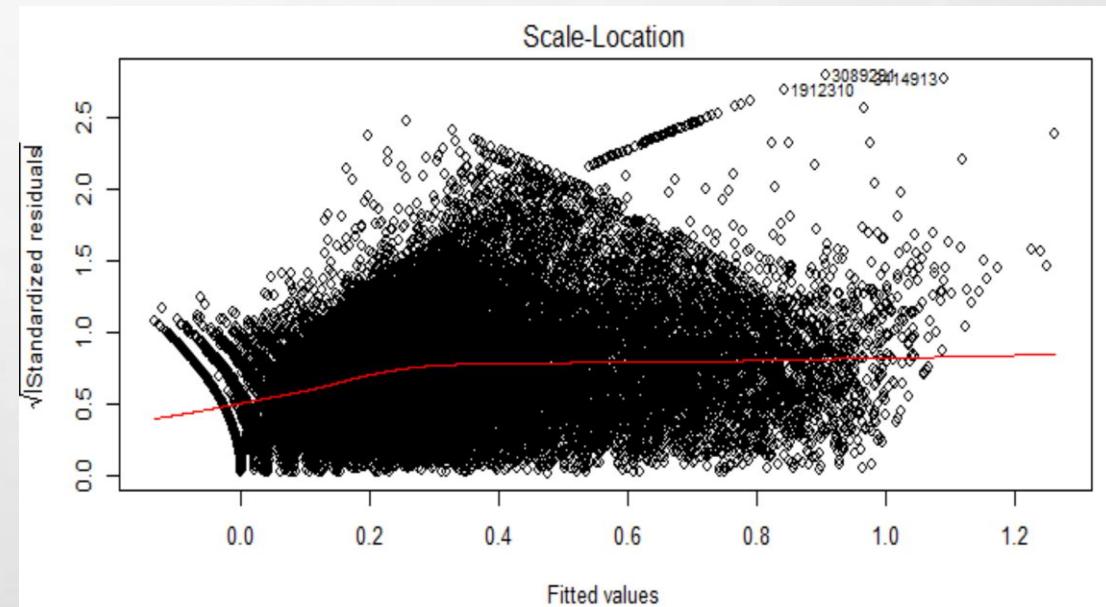
## ASSUMPTIONS

- **INDEPENDENCE**

All predictors must be independent of each other.

- **HOMOSCEDASTICITY**

Plot of standardized residuals versus predicted values can show whether points are equally distributed across all values of the independent variables.



# MULTIPLE LINEAR REGRESSION

## RESULTS

```
YHAT = PREDICT(OBJ_FWD, NEWDATA=TEST_DF[,])
```

```
YTEST = TEST_DF$WINPLACEPERC
```

```
RMSE(YTEST, YHAT)
```

```
[1] 0.1163719 → RMSE
```

### Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	0.3040876	0.0006420	473.631	< 2e-16 ***
walkDistance	0.0971712	0.0009572	101.516	< 2e-16 ***
killPlace	-0.2096509	0.0014117	-148.513	< 2e-16 ***
killstreaks	-0.0807176	0.0019531	-41.328	< 2e-16 ***
numGroups	0.0302902	0.0006503	46.580	< 2e-16 ***
matchDuration	-0.0291164	0.0006517	-44.677	< 2e-16 ***
kills	-0.0631869	0.0020713	-30.506	< 2e-16 ***
boosts	0.0122526	0.0008523	14.376	< 2e-16 ***
weaponsAcquired	0.0150136	0.0008272	18.150	< 2e-16 ***
assists	0.0060512	0.0006700	9.032	< 2e-16 ***
DBNOS	-0.0105053	0.0010303	-10.196	< 2e-16 ***
heals	0.0066036	0.0007602	8.687	< 2e-16 ***
revives	0.0055498	0.0006629	8.372	< 2e-16 ***
killPoints	0.0034375	0.0006465	5.317	1.06e-07 ***
longestKill	0.0052059	0.0009684	5.376	7.68e-08 ***
headshotKills	-0.0036143	0.0007365	-4.907	9.28e-07 ***
damageDealt	0.0031530	0.0010672	2.954	0.00313 **
swimDistance	0.0015185	0.0006192	2.452	0.01420 *

---

Signif. codes: 0 ‘\*\*\*’ 0.001 ‘\*\*’ 0.01 ‘\*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Residual standard error: 0.1165 on 32912 degrees of freedom

Multiple R-squared: 0.7891, Adjusted R-squared: 0.7889

F-statistic: 7242 on 17 and 32912 DF, p-value: < 2.2e-16

# REGRESSION TREE

## Predicted Rating

If killPlace value lies between 0.66 and 0.99, then the winning percentage will be 0.13

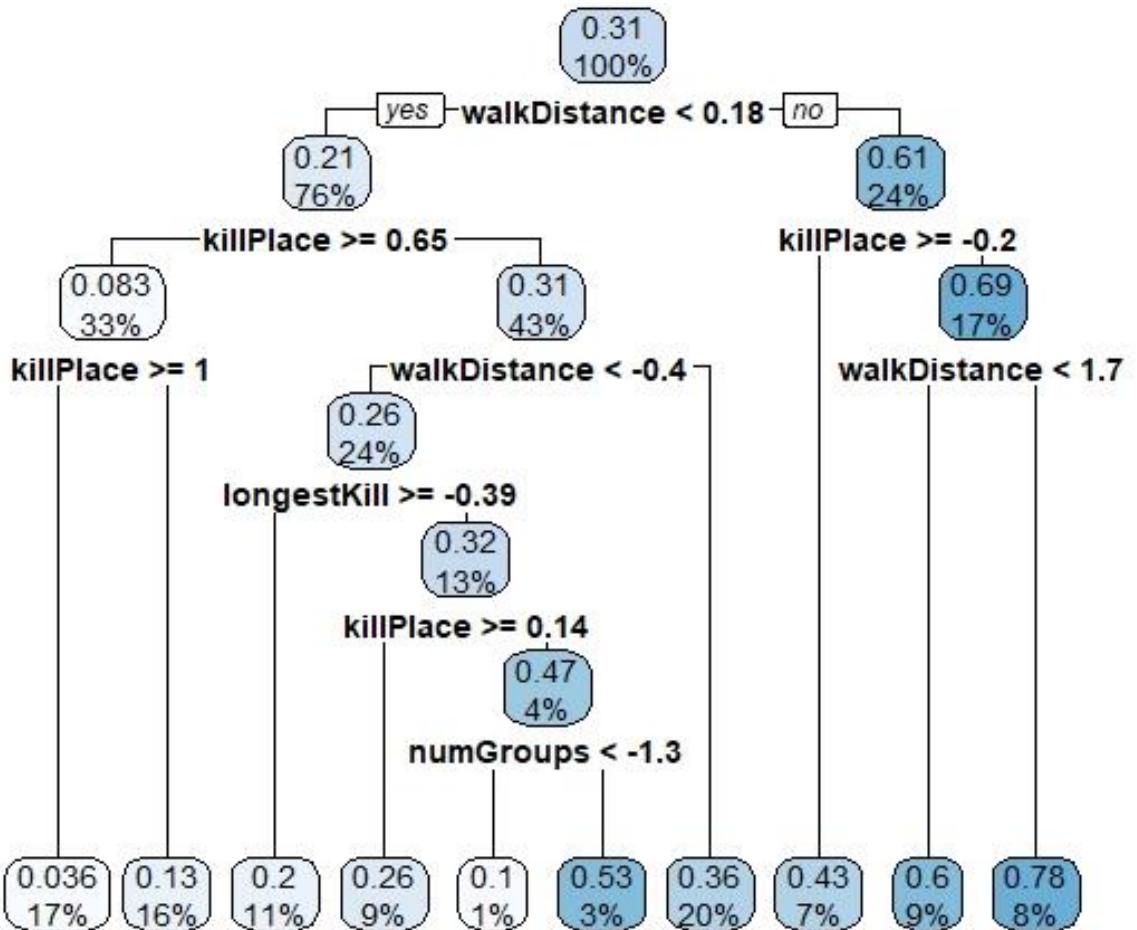
## Decision nodes with variable importance

- Walkdistance: 43
- Killplace: 20
- Longestkill: 4
- Numgroups: 1

# REGRESSION TREE

Best Pruned Tree

RMSE: 0.1278694



# REGRESSION TREE

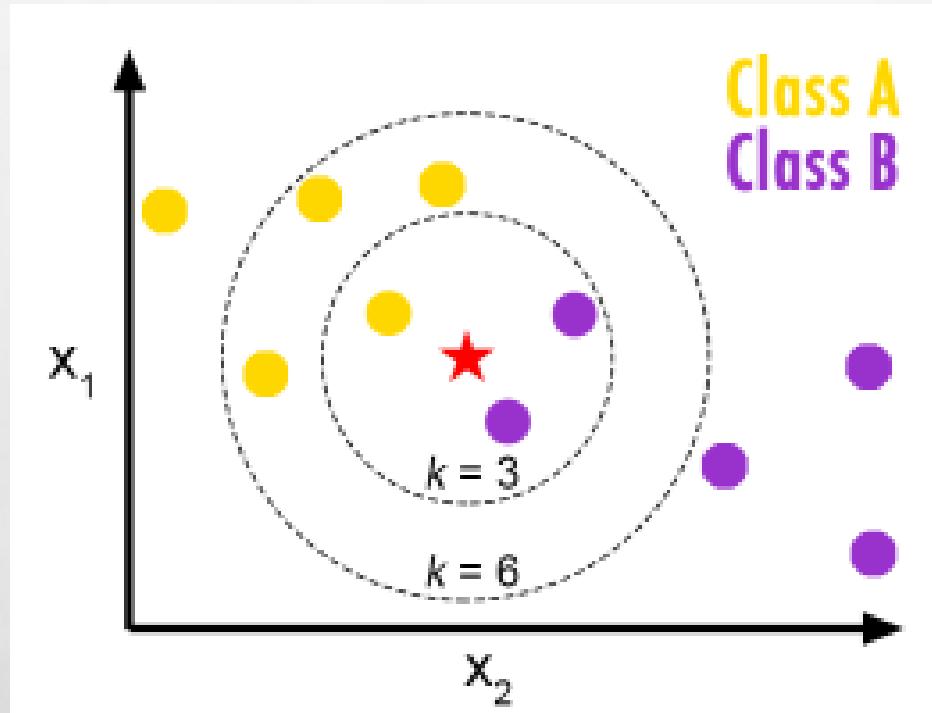
```
> rpart.plot(cart_fit)
> pfit = prune(cart_fit, cp = cart_fit$cptable[which.min(cart_fit$cptable[, "xerror"]),"CP"])
> rpart.plot(pfit)
```

Output:

```
> rmse(yhat.test, ytest)
[1] 0.1278694
> rmse(yhat.test_prune, ytest)
[1] 0.1278694
```

CP	nsplit	rel error	xerror	xstd
1 0.458263	0	1 1.000063	0.007357	
2 0.145269	1 0.541737	0.5439	0.005379	
3 0.052109	2 0.396468	0.398624	0.004279	
4 0.021327	3 0.344359	0.346266	0.004052	
5 0.016359	4 0.323032	0.32504	0.003942	
6 0.012975	7 0.273955	0.276986	0.003498	
7 0.012213	8 0.260981	0.272975	0.003467	
8 0.01	9 0.248768	0.25355	0.003518	

# KNN – K NEAREST NEIGHBORS



# KNN – K NEAREST NEIGHBORS

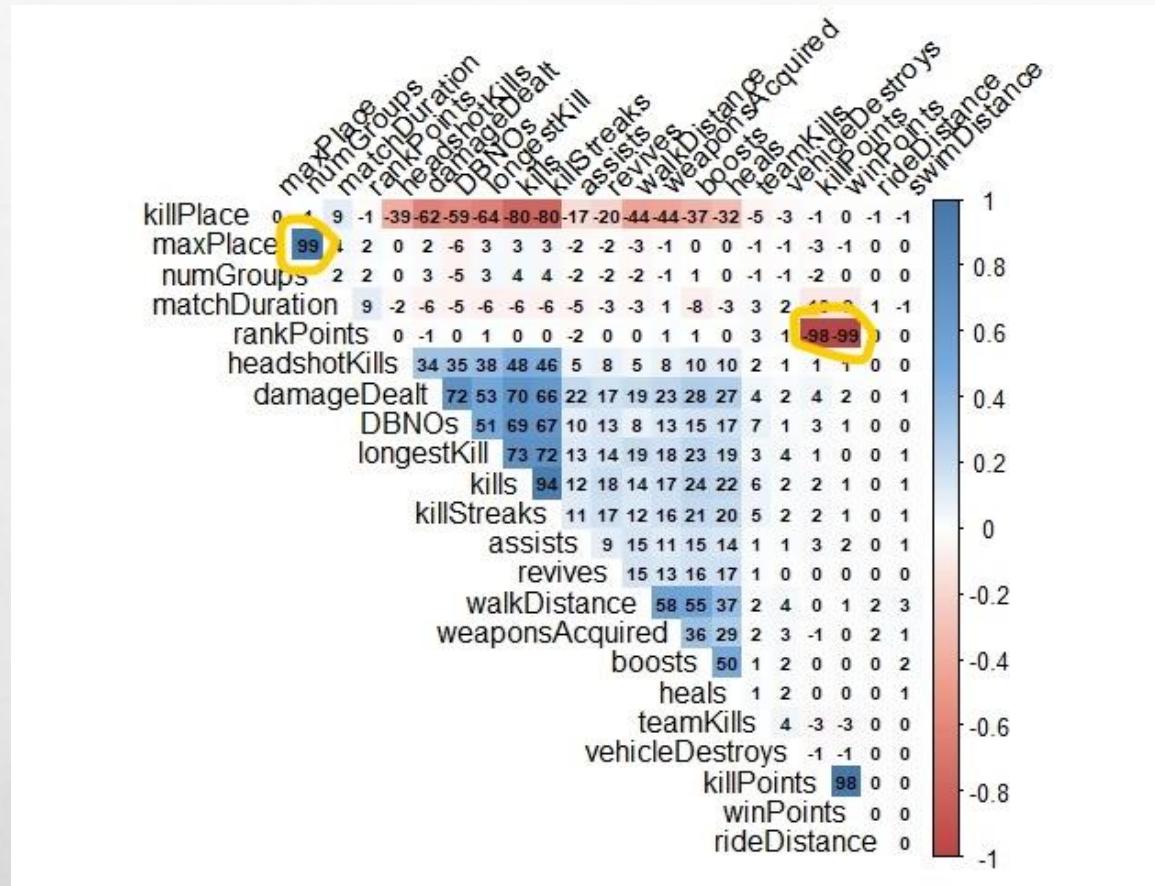
## Optimum K

```
> knn.reg.bestK = function(Xtrain, Xtest, ytrain, ytest, kmax=NULL) {  
+   if (is.null(kmax)) kmax = length(ytrain)^.5  
+   vec.rmse = rep(NA, kmax)  
+   for (k in 1:kmax) {  
+     yhat.test = knn.reg(Xtrain, Xtest, ytrain, k)$pred  
+     vec.rmse[k] = rmse(yhat.test, ytest)  
+   }  
+   list(k.opt = which.min(vec.rmse), rmse.min = min(vec.rmse), vec.rmse)  
+ }  
> knn.reg.bestK(train_df[,1:20], test_df[,1:20], train_df[,21], test_df[,21])  
$`k.opt'  
[1] 11  
  
$rmse.min  
[1] 0.1164329  
  
[[3]]  
[1] 0.1489087 0.1292303 0.1235195 0.1208125 0.1189639 0.1178877 0.1173630 0.1167980 0.1164932  
[10] 0.1165901 0.1164329 0.1165463 0.1166584 0.1166734 0.1167291 0.1169729 0.1169507 0.1171207  
[19] 0.1172004 0.1173454 0.1174112 0.1175510 0.1175958 0.1176799 0.1178767 0.1180416 0.1181332  
[28] 0.1183772 0.1185622 0.1187579 0.1189284 0.1191680 0.1193679 0.1194682 0.1196381 0.1197152  
[37] 0.1199071 0.1200618 0.1202195 0.1203314 0.1204788 0.1206817 0.1208332 0.1209504 0.1210391
```

# KNN – K NEAREST NEIGHBORS

```
> knn_obj <- kknn(winPlacePerc~, train_df, test_df, k=11)
> names(knn_obj)
[1] "fitted.values" "CL"          "w"           "D"           "c"
[6] "prob"          "response"    "distance"    "call"        "terms"
> head(knn_obj$D[1:2,1:11])
     [,1]   [,2]   [,3]   [,4]   [,5]   [,6]   [,7]   [,8]   [,9]
[1,] 1.7706986 1.7731392 1.8984154 1.9963801 2.0224436 2.0795305 2.0921001 2.1677385 2.1833689
[2,] 0.4299748 0.4889075 0.5464724 0.5548825 0.5832166 0.6005727 0.6326097 0.6489151 0.6865022
     [,10]  [,11]
[1,] 2.2006689 2.2373768
[2,] 0.7151994 0.7310168
> head(knn_obj$fitted.values)
[1] 0.287107873 0.151477529 0.149663675 0.000000000 0.008124091 0.216216771
> head(knn_obj$D[1:2,1:11])
     [,1]   [,2]   [,3]   [,4]   [,5]   [,6]   [,7]   [,8]   [,9]
[1,] 1.7706986 1.7731392 1.8984154 1.9963801 2.0224436 2.0795305 2.0921001 2.1677385 2.1833689
[2,] 0.4299748 0.4889075 0.5464724 0.5548825 0.5832166 0.6005727 0.6326097 0.6489151 0.6865022
     [,10]  [,11]
[1,] 2.2006689 2.2373768
[2,] 0.7151994 0.7310168
> head(knn_obj$response)
[1] "continuous"
> head(knn_obj$call)
kknn(formula = winPlacePerc ~ ., train = train_df, test = test_df,
      k = 11)
> knn_obj$terms
winPlacePerc ~ assists + boosts + damageDealt + DBNOS + headshotkills +
  heals + killPlace + killPoints + kills + killStreaks + longestKill +
  matchDuration + numGroups + revives + rideDistance + swimDistance +
  teamkills + vehicleDestroys + walkDistance + weaponsAcquired
```

## kNN objects



# ANALYSIS

- Applied PCA on kill related variables.

```
pca_obj <- prcomp(train_df[,c(5,7,9,10,11)])
```

# RESULTS

## MLR OUTPUT WITH PCA

```
Residual standard error: 0.1485 on 32917 degrees of freedom
Multiple R-squared:  0.6574 ,   Adjusted R-squared: 0.6573
F-statistic:  5264 on 12 and 32917 DF,  p-value: < 2.2e-16
```

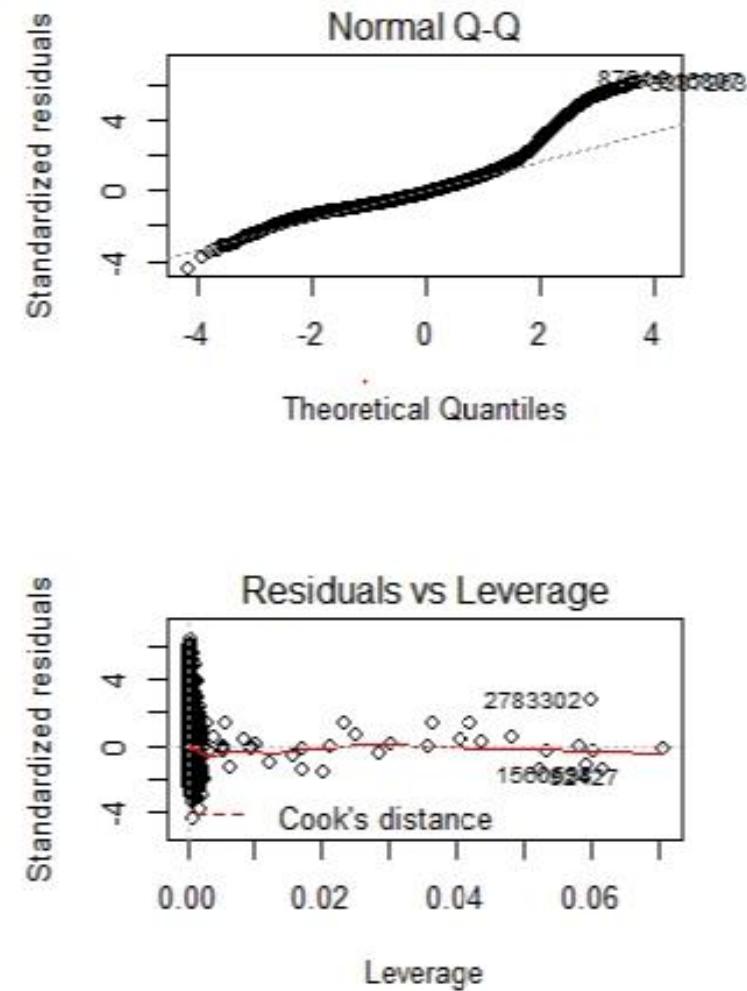
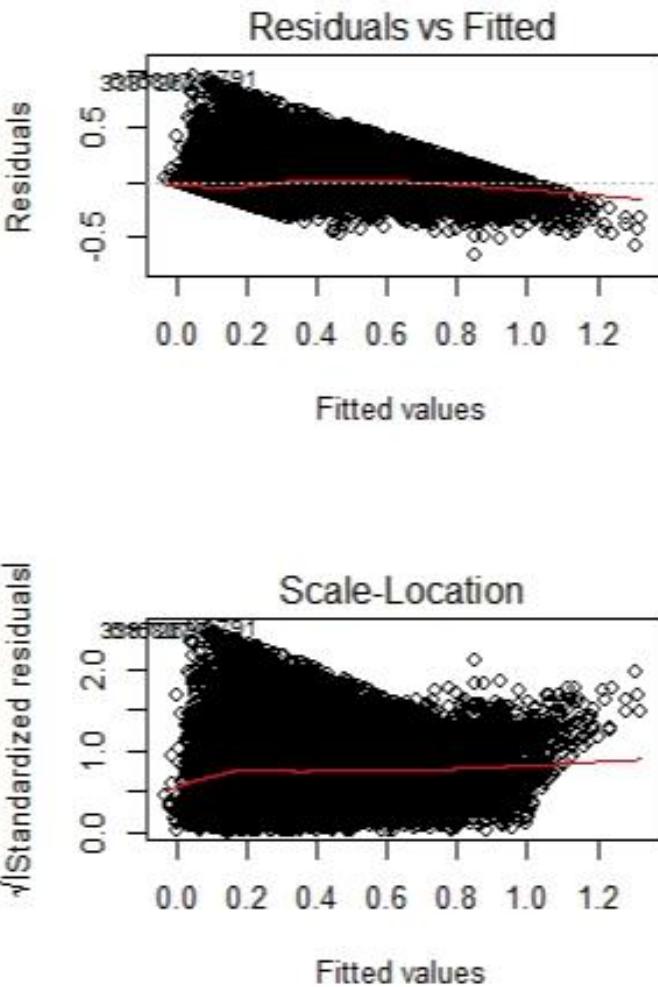
## MLR OUTPUT WITHOUT PCA

```
Residual standard error: 0.1178 on 33183 degrees of freedom
Multiple R-squared:  0.7849 ,   Adjusted R-squared: 0.7848
F-statistic:  7121 on 17 and 33183 DF,  p-value: < 2.2e-16
```

# RESULTS

## MLR assumptions with PCA

```
par(mfrow=c(2,2))  
plot(obj_fwd)
```



# RESULTS

	RMSE	
	Without PCA	With PCA
MLR	0.1162	0.1479
CART	0.1256	0.1273
k-NN	0.1160	0.1319

# CONCLUSION



Analysis without applying PCA provides the best results for all the three algorithms.



MLR is not an optimal method to predict win place percentile as the assumptions of MLR does not meet in either of the case.



Based on the RMSE value k-NN with  $k = 11$  is the best method to predict the win place percentile variable

# PREDICTED OUTPUT

```
> head(df_output,34)
   Id      winPlacePerc
1 4125d337a9bc41 0.0210994094581172
2 8a35b912403f72 0.118658789158246
3 209943df673589 0.452947508854676
4 011ce2c63edda6 0.47935400822543
5 e865b7c0755bbe 0.0865101575176968
6 31faec823f564b 0.760342888276353
7 d1aed2cd37df1f 0.712297257821613
8 e9c9d620252c77 0.106515898862378
9 4fa2e0712efa15 0.282700405389399
10 93234b66565eeef 0.241654342100448
11 c72dba44deeae7 0.580401924676827
12 2f30b13bf09bba 0.148551242971799
13 33fab8d8035f58 0.0700019494104264
14 0b2276e56128ef 0.33655708136644
15 34990f1d2378bd 0.239701090601804
16 3aa86db23adf23 0.077701295739267
17 4f22ab218e3ebd 0.234635477001815
18 51324ce6157f61 0.258170785297725
19 ad8e72c55070b1 0.469025704317451
20 5bd89dd6b03179 0.201024376983084
21 a8f0e6dd80cdac 0.477613045647615
22 fe2a790031dd15 0.385390098095539
23 31adbd1645e520 0.126201990626131
24 bc4d5b63a0019e 0.0640082342752
25 4a5b45bc5c0324 0.57950009519022
26 db4289b05a135f 0.489947786657775
27 40a90775dc0c0b 0.298227135327823
28 e3057bbf900782 0.320867812608204
29 295641fbbe6bfd 0.0231543226945885
30 f302dde0e3e487 0.39531421355708
31 386ef41cba8f09 0.267115585869294
32 f1f080e1c2f3e9 0.48324077641656
33 22ff8b3fc32fdc 0.197296916803907
34 c128d433b0e6c8 0.396629780646545
```

# BEST STRATEGY

